


目录

DDPM完全解读
名词解析
预定义
Diffusion Process
Reverse Process
Loss
推导
$q(x_T x_0)$
$q(x_{t-1} x_t, x_0)$
$p_\theta(x_{t-1} x_t)$
首先考虑std
其次考虑mean
整合
Reference

DDPM完全解读

名词解析



- diffusion process: also called forward process, training process, represented by $q(x_t)$
- reverse process: also called sampling process, inference process, represented by $p(x_t)$

注意:

- diffusion process: is fixed to a Markov chain that gradually adds Gaussian noise to the data according to variance schedule β_1, \dots, β_T , 即变换前后满足高斯分布, 当前状态只与前一时刻有关; 下一小节将给出diffusion过程的分布预定义形式, 即 **variance schedule** 是自定义的constant; 且从上图可以看出diffusion过程与 θ 无关, 只是为了求loss, 将 θ 作用在forward input上, 后续将具体介绍
- reverse process: is defined as a Markov chain with learned Gaussian transition starting at $P(x_T) := N(x_T; 0, I)$, 但reverse过程的mean and std是与 θ 相关的函数, 为了使 $p_\theta(x_0)$ 尽量接近 $q_{\text{data}}(x_0)$, 需要找到mean和std的最佳定义, 使likelihood of $p_\theta(x_0)$ 最大, 这也是diffusion model的loss定义, 后续将具体介绍

预定义

Diffusion Process

根据Diffusion Model[1]的定义, 定义了diffusion process is fixed to a Markov chain that gradually adds Gaussian noise to the data according to variance schedule β_1, \dots, β_T :

$$q(x_t|x_{t-1}) := N(x_t; \sqrt{1 - \beta_t}x_{t-1}, \beta_t I)$$
$$q(x_{1:T}|x_0) := \prod_{t=1}^T q(x_t|x_{t-1})$$

where β_t is variance schedule, also called diffusion rate.

注意此处 $[\cdot]$ 不是阶乘，代表嵌套的意思，类似normalizing flow里的表达

Reverse Process

根据定义，sampling/reverse process is defined as a Markov chain with learned Gaussian transition starting at $p(x_T) := N(x_T; 0, I)$:

$$\begin{aligned} p(x_T) &:= N(x_T; 0, I) \\ p_\theta(x_{t-1}|x_t) &:= N(x_{t-1}; \mu_\theta(x_t, t), \sum_\theta(x_t, t)I) \\ p_\theta(x_{0:T}) &= p_\theta(x_0, x_1, \dots, x_T) = p(x_T) \cdot \prod_{t=1}^T p_\theta(x_{t-1}|x_t) \end{aligned}$$

注意此处 $[\cdot]$ 不是阶乘，代表嵌套的意思，类似normalizing flow里的表达

Loss

Reverse process is θ parameterization, 为了使reverse process尽可能的得到fidelity result, 则需要找到使 $p_\theta(x_0)$ 最接近 $q_{data}(x_0)$ 分布的参数 θ , 即采用maximize log likelihood estimation, 等价于minimize negative log likelihood.

下式中的 $p_\theta(x_0)$ 通常很难准确求解, Sampling process start from $p_{latent}(x_T) := N(x_T; 0, I)$, 则:

$$\begin{aligned} p_\theta(x_0, x_1, \dots, x_{T-1}|x_T) \cdot p(x_T) &= p_\theta(x_0, x_1, \dots, x_T) \\ p_\theta(x_0) &= \int p_\theta(x_0, x_1, \dots, x_T) dx_1 dx_2 \dots dx_T = \int p_\theta(x_{0:T}) dx_{1:T} \end{aligned}$$

故 θ 通过minimize variational bound on log likelihood求解, log likelihood定义为:

$$\begin{aligned} \mathbb{E}_{q_{data}(x_0)} - \log p_\theta(x_0) &= -\mathbb{E}_{q_{data}(x_0)} \log \mathbb{E}_{q(x_1, \dots, x_T|x_0)} \left[\frac{p_\theta(x_0, x_1, \dots, x_{T-1}|x_T) \cdot p(x_T)}{q(x_1, \dots, x_T|x_0)} \right] \\ &\text{计算其variational bound, 将}\mathbb{E}\text{都提前, 并合并下标概率, 可得:} \\ \mathbb{E}_{q_{data}(x_0)} - \log p_\theta(x_0) &\leq -\mathbb{E}_{q(x_0, \dots, x_T)} \log \left[\frac{p_\theta(x_0, x_1, \dots, x_{T-1}|x_T) \cdot p(x_T)}{q(x_1, \dots, x_T|x_0)} \right] \\ &\text{简写为} \\ \mathbb{E}_{q_{data}(x_0)} - \log p_\theta(x_0) &\leq -\mathbb{E}_{q(x_{0:T})} \log \left[\frac{p_\theta(x_{0:T})}{q(x_{1:T}|x_0)} \right] =: L \end{aligned}$$

将上式 L 进一步可简化为:

$$\mathbb{E}_q \left[\underbrace{D_{KL}(q(\mathbf{x}_T|x_0) \parallel p(\mathbf{x}_T))}_{L_T} + \sum_{t>1} \underbrace{D_{KL}(q(\mathbf{x}_{t-1}|x_t, x_0) \parallel p_\theta(\mathbf{x}_{t-1}|x_t))}_{L_{t-1}} - \underbrace{\log p_\theta(x_0|x_1)}_{L_0} \right] \quad (5)$$

具体推导如下 (注以下推导没加负号[2], 加负号的推导见[3]-Extra information-A Extended derivations) :

Proof. We expand the ELBO in Eq. (3) into the sum of a sequence of tractable KL divergences below.

$$\begin{aligned} \text{ELBO} &= \mathbb{E}_q \log \frac{p_\theta(x_0, \dots, x_{T-1}|x_T) \times p_{latent}(x_T)}{q(x_1, \dots, x_T|x_0)} \\ &= \mathbb{E}_q \left(\log p_{latent}(x_T) - \sum_{t=1}^T \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_t|x_{t-1})} \right) \\ &= \mathbb{E}_q \left(\log p_{latent}(x_T) - \log \frac{p_\theta(x_0|x_1)}{q(x_1|x_0)} - \sum_{t=2}^T \left(\log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} + \log \frac{q(x_{t-1}|x_0)}{q(x_t|x_0)} \right) \right) \\ &= \mathbb{E}_q \left(\log \frac{p_{latent}(x_T)}{q(x_T|x_0)} - \log p_\theta(x_0|x_1) - \sum_{t=2}^T \log \frac{p_\theta(x_{t-1}|x_t)}{q(x_{t-1}|x_t, x_0)} \right) \\ &= -\mathbb{E}_q \left(\text{KL}(q(x_T|x_0) \parallel p_{latent}(x_T)) + \sum_{t=2}^T \text{KL}(q(x_{t-1}|x_t, x_0) \parallel p_\theta(x_{t-1}|x_t)) - \log p_\theta(x_0|x_1) \right) \end{aligned} \quad (9)$$

注意, loss的计算需要用到

$$q(x_T|x_0)$$

$$q(x_{t-1}|x_t, x_0)$$

$$p_\theta(x_{t-1}|x_t), p_\theta(x_0|x_1)$$

前两个分布以下将分别推导，第三个分布即DDPM[3] proposed distribution, which makes DDPM resembling denoising score matching[4].

推导

$q(x_T|x_0)$

根据diffusion model预定义，从1至T的任意时刻， x_t 相对 x_0 的后验： $q(x_{1:T}|x_0) := \prod_{t=1}^T q(x_t|x_{t-1})$

以 $q(x_3|x_0)$ 为例，即 $t = 3$ 时：

$$q(x_1|x_0) := N(x_1; \sqrt{1-\beta_1}x_0, \beta_1 I_1), \text{ so } x_1 = \sqrt{1-\beta_1}x_0 + \sqrt{\beta_1}I_1$$

$$q(x_2|x_1) := N(x_2; \sqrt{1-\beta_2}x_1, \beta_2 I_2), \text{ so } x_2 = \sqrt{1-\beta_2}(\sqrt{1-\beta_1}x_0 + \sqrt{\beta_1}I_1) + \sqrt{\beta_2}I_2$$

$$q(x_3|x_2) := N(x_3; \sqrt{1-\beta_3}x_2, \beta_3 I_3), \text{ so } x_3 = \sqrt{1-\beta_3}[\sqrt{1-\beta_2}(\sqrt{1-\beta_1}x_0 + \sqrt{\beta_1}I_1) + \sqrt{\beta_2}I_2] + \sqrt{\beta_3}I_3$$

令 $\alpha_t = 1 - \beta_t, \bar{\alpha}_t = \prod_{s=1}^t \alpha_s$ ，则 $q(x_3|x_0)$ 的均值，方差为

$$\text{mean} = \sqrt{1-\beta_3}\sqrt{1-\beta_2}\sqrt{1-\beta_1} = \sqrt{\bar{\alpha}_3}$$

$$\text{std}^2 = (\sqrt{1-\beta_3}\sqrt{1-\beta_2}\sqrt{\beta_1})^2 + (\sqrt{1-\beta_3}\sqrt{\beta_2})^2 + (\sqrt{\beta_3})^2 = \alpha_3\alpha_2(1-\alpha_1) + \alpha_3(1-\alpha_2) + (1-\alpha_3) = 1 - \alpha_1\alpha_2\alpha_3 = 1 - \bar{\alpha}_3$$

因此， $q(x_3|x_0) := N(x_3; \sqrt{\bar{\alpha}_3} \cdot x_0, (1 - \bar{\alpha}_3)I)$ ，同理，推广到所有的t可得：

$$q(x_t|x_0) := N(x_t; \sqrt{\bar{\alpha}_t} \cdot x_0, (1 - \bar{\alpha}_t)I)$$

或者，迭代的理论推导如下：

$$\begin{aligned} x_t &= \sqrt{\alpha_t}x_{t-1} + \sqrt{\beta_t}\epsilon_t \\ &= \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{\alpha_t\beta_{t-1}}\epsilon_{t-1} + \sqrt{\beta_t}\epsilon_t \\ &= \sqrt{\alpha_t\alpha_{t-1}\alpha_{t-2}}x_{t-3} + \sqrt{\alpha_t\alpha_{t-1}\beta_{t-2}}\epsilon_{t-2} + \sqrt{\alpha_t\beta_{t-1}}\epsilon_{t-1} + \sqrt{\beta_t}\epsilon_t \\ &= \dots \\ &= \sqrt{\bar{\alpha}_t}x_0 + \sqrt{\alpha_t\alpha_{t-1}\dots\alpha_2\beta_1}\epsilon_1 + \dots + \sqrt{\alpha_t\beta_{t-1}}\epsilon_{t-1} + \sqrt{\beta_t}\epsilon_t \end{aligned}$$

Note that $q(x_t|x_0)$ is still Gaussian, and the mean of x_t is $\sqrt{\bar{\alpha}_t}x_0$, and the variance matrix is $(\alpha_t\alpha_{t-1}\dots\alpha_2\beta_1 + \dots + \alpha_t\beta_{t-1} + \beta_t)I = (1 - \bar{\alpha}_t)I$. Therefore,

$$q(x_t|x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I). \quad (10)$$

It is worth mentioning that,

$$q(x_T|x_0) = \mathcal{N}(x_T; \sqrt{\bar{\alpha}_T}x_0, (1 - \bar{\alpha}_T)I), \quad (11)$$

where $\bar{\alpha}_T = \prod_{t=1}^T (1 - \beta_t)$ approaches zero with large T .

$q(x_{t-1}|x_t, x_0)$

根据贝叶斯公式，全概率公式等，可得：

$$\begin{aligned} q(x_{t-1}|x_t, x_0) &= \frac{q(x_{t-1}, x_t, x_0)}{q(x_t, x_0)} = \frac{q(x_t|x_{t-1}, x_0) \cdot q(x_0, x_{t-1})}{q(x_t, x_0)} = \frac{q(x_t|x_{t-1}, x_0) \cdot q(x_0, x_{t-1})}{q(x_t, x_0)} \\ &= \frac{q(x_t|x_{t-1}, x_0) \cdot q(x_{t-1}|x_0) \cdot q(x_0)}{q(x_t|x_0) \cdot q(x_0)} = \frac{q(x_t|x_{t-1}, x_0) \cdot q(x_{t-1}|x_0)}{q(x_t|x_0)} \end{aligned}$$

根据各维独立高斯分布，将 x_{t-1} 关于 x_t, x_0 的分布可进一步化简得：

Next, by Bayes rule and Markov chain property,

$$\begin{aligned}
q(x_{t-1}|x_t, x_0) &= \frac{q(x_t|x_{t-1}) q(x_{t-1}|x_0)}{q(x_t|x_0)} \\
&= \frac{\mathcal{N}(x_t; \sqrt{\alpha_t}x_{t-1}, \beta_t I) \mathcal{N}(x_{t-1}; \sqrt{\bar{\alpha}_{t-1}}x_0, (1 - \bar{\alpha}_{t-1})I)}{\mathcal{N}(x_t; \sqrt{\bar{\alpha}_t}x_0, (1 - \bar{\alpha}_t)I)} \\
&= (2\pi\beta_t)^{-\frac{d}{2}} (2\pi(1 - \bar{\alpha}_{t-1}))^{-\frac{d}{2}} (2\pi(1 - \bar{\alpha}_t))^{-\frac{d}{2}} \times \\
&\quad \exp\left(-\frac{\|x_t - \sqrt{\alpha_t}x_{t-1}\|^2}{2\beta_t} - \frac{\|x_{t-1} - \sqrt{\bar{\alpha}_{t-1}}x_0\|^2}{2(1 - \bar{\alpha}_{t-1})} + \frac{\|x_t - \sqrt{\bar{\alpha}_t}x_0\|^2}{2(1 - \bar{\alpha}_t)}\right) \\
&= (2\pi\tilde{\beta}_t)^{-\frac{d}{2}} \exp\left(-\frac{1}{2\tilde{\beta}_t} \left\|x_{t-1} - \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0 - \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t\right\|^2\right)
\end{aligned}$$

Therefore,

$$q(x_{t-1}|x_t, x_0) = \mathcal{N}(x_{t-1}; \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0 + \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t, \tilde{\beta}_t I). \quad (12)$$

其中, $\tilde{\beta}_t$ 代表:

$$\tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \cdot \beta_t$$

$p_\theta(x_{t-1}|x_t)$

由于 L_{t-1} 是 $q(x_{t-1}|x_t, x_0)$ 与 $p_\theta(x_{t-1}|x_t)$ 的KL散度, $q(x_{t-1}|x_t, x_0)$ 已求得, 因此需要确定 $p_\theta(x_{t-1}|x_t)$ 的分布, 使loss有closed-form calculation, 因此DDPM[3]给出了 $p_\theta(x_{t-1}|x_t) := N(x_{t-1}; \mu_\theta(x_t, t), \sum_\theta(x_t, t)I)$ 中 $\mu_\theta(x_t, t), \sum_\theta(x_t, t)$ 的形式, 实现了closed-form expression, 并与denoising score matching对应起来, 以下将具体介绍。

根据前文reverse process预定义,

$$p_\theta(x_{t-1}|x_t) := N(x_{t-1}; \mu_\theta(x_t, t), \sum_\theta(x_t, t)I)$$

Loss中与其相关的项为:

$$L_{t-1} = \sum_{t>1} D_{KL}(q(x_{t-1}|x_t, x_0) || p_\theta(x_{t-1}|x_t))$$

其中, 根据上文推导, 将 $q(x_{t-1}|x_t, x_0)$ 简写为:

$$\begin{aligned}
q(x_{t-1}|x_t, x_0) &:= N(x_{t-1}; \bar{\mu}_t(x_t, x_0), \bar{\beta}_t I) \\
\text{where } \bar{\mu}_t(x_t, x_0) &= \frac{\sqrt{\bar{\alpha}_{t-1}} \cdot \beta_t}{1 - \bar{\alpha}_t} \cdot x_0 + \frac{\sqrt{\alpha_t} \cdot (1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} \cdot x_t
\end{aligned}$$

首先考虑std

Ho通过实验发现, 令

- $\sum_\theta(x_t, t) = \sigma_t^2 = \beta_t, \beta_t \text{趋向} 1$
- $\sum_\theta(x_t, t) = \sigma_t^2 = \bar{\beta}_t, \bar{\beta}_t < \beta_t$

有相似的结果, The first choice is optimal for $x_0 := N(x_0; 0, I)$, and the second is optimal for x_0 deterministically set to one point. These are the two extreme choices corresponding to upper and lower bounds on reverse process entropy for data with coordinatewise unit variance[1].

无论第一种还是第二种方式, $\sum_\theta(x_t, t)$ 都与 θ 无关, 由于 $q(x_{t-1}|x_t, x_0)$ 的variance与 θ 也无关, 因此两个分布的std项带入KL divergence中计算得到常数C

因此, 实验中采用第二种方式。

其次考虑mean

将C带入 L_{t-1} 可化简得:

$$L_{t-1} = \mathbb{E}_q \left[\frac{1}{2\sigma_t^2} \|\tilde{\mu}_t(x_t, x_0) - \mu_\theta(x_t, t)\|^2 \right] + C \quad (8)$$

根据上文:

$$q(x_t|x_0) := N(x_t; \sqrt{\alpha_t} \cdot x_0, (1 - \alpha_t)I) \rightarrow x_t = \sqrt{\alpha_t} \cdot x_0 + \sqrt{1 - \alpha_t} \cdot \epsilon \rightarrow x_0 = \frac{1}{\sqrt{\alpha_t}} \cdot (x_t - \sqrt{1 - \alpha_t} \cdot \epsilon)$$

$$\bar{\mu}_t(x_t, x_0) = \frac{\sqrt{\alpha_{t-1}} \cdot \beta_t}{1 - \alpha_t} \cdot x_0 + \frac{\sqrt{\alpha_t} \cdot (1 - \alpha_{t-1})}{1 - \alpha_t} \cdot x_t$$

带入上式得:

$$L_{t-1} - C = \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[\frac{1}{2\sigma_t^2} \left\| \tilde{\mu}_t \left(\mathbf{x}_t(\mathbf{x}_0, \epsilon), \frac{1}{\sqrt{\alpha_t}} (\mathbf{x}_t(\mathbf{x}_0, \epsilon) - \sqrt{1 - \alpha_t} \epsilon) \right) - \mu_\theta(\mathbf{x}_t(\mathbf{x}_0, \epsilon), t) \right\|^2 \right] \quad (9)$$

$$= \mathbb{E}_{\mathbf{x}_0, \epsilon} \left[\frac{1}{2\sigma_t^2} \left\| \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t(\mathbf{x}_0, \epsilon) - \frac{\beta_t}{\sqrt{1 - \alpha_t}} \epsilon \right) - \mu_\theta(\mathbf{x}_t(\mathbf{x}_0, \epsilon), t) \right\|^2 \right] \quad (10)$$

根据上式可知, μ_θ 需要预测 (尽可能近似) $\frac{1}{\sqrt{\alpha_t}} (x_t(x_0, \epsilon) - \frac{\beta_t}{\sqrt{1 - \alpha_t}} \epsilon)$, 由于其中 $x_t(x_0, \epsilon)$ 是reverse input, 是已知的 (gaussian noise), 因此, μ_θ 的proposition只能将参数作用于 ϵ , 同时将 x_t 也作为其输入, 因此Ho将mean定义为:

$$\mu_\theta(\mathbf{x}_t, t) = \tilde{\mu}_t \left(\mathbf{x}_t, \frac{1}{\sqrt{\alpha_t}} (\mathbf{x}_t - \sqrt{1 - \alpha_t} \epsilon_\theta(\mathbf{x}_t)) \right) = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{\beta_t}{\sqrt{1 - \alpha_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) \quad (11)$$

其中, ϵ_θ is a function approximator,用来根据 x_t 估计 ϵ ,即gaussian noise.

整合

有了std和mean的定义, 即可给出 x_{t-1} 的解析:

$$x_{t-1} = \text{mean} + \text{std} \cdot z = \frac{1}{\sqrt{\alpha_t}} \cdot \left(x_t - \frac{\beta_t}{\sqrt{1 - \alpha_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t z$$

where $z := N(0, I)$

x_{t-1} 既可以视为reverse所得, 也可视为diffusion所得, 因此training过程, ϵ_θ 可作用于diffusion variable

The complete sampling procedure resembles Langevin dynamics with ϵ_θ as a learned gradient of the data density. 同时, Eq10.可简化为:

$$\mathbb{E}_{x_0, \epsilon} \left[\frac{\beta_t^2}{2\sigma_t^2 \alpha_t (1 - \alpha_t)} \left\| \epsilon - \epsilon_\theta(\sqrt{\alpha_t} \cdot x_0 + \sqrt{1 - \alpha_t} \cdot \epsilon, t) \right\|^2 \right]$$

which resembles denoising score matching over multiple noise scales indexed by t. 上式 is equal to (one term of) the variational bound for the Langevin-like reverse process, we see

that optimizing an objective resembling denoising score matching is equivalent to using variational inference to fit the finite-time marginal of a sampling chain resembling Langevin dynamics. 即与denoising score matching对应起来了。

此外, 通过不同的化简形式, 也可以约掉 x_t , 使 μ_θ 化简为关于 x_0 , 但作者通过实验发现这种方式生成的图片质量更差。

因此, DDPM的diffusion process和reverse process可概括如下:

Algorithm 1 Training

- 1: **repeat**
- 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
- 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
- 4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 5: Take gradient descent step on $\nabla_\theta \left\| \epsilon - \epsilon_\theta(\sqrt{\alpha_t} \mathbf{x}_0 + \sqrt{1 - \alpha_t} \epsilon, t) \right\|^2$
- 6: **until** converged

Algorithm 2 Sampling

- 1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for** $t = T, \dots, 1$ **do**
- 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
- 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \alpha_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
- 5: **end for**
- 6: **return** \mathbf{x}_0

Reference

- [1]. Deep Unsupervised Learning using Nonequilibrium Thermodynamics, 2015.
- [2]. DIFFWAVE: A VERSATILE DIFFUSION MODEL FOR AUDIO SYNTHESIS, 2021.
- [3]. Denoising Diffusion Probabilistic Models, 2020.
- [4]. Generative Modeling by Estimating Gradients of the Data Distribution, 2019.

